

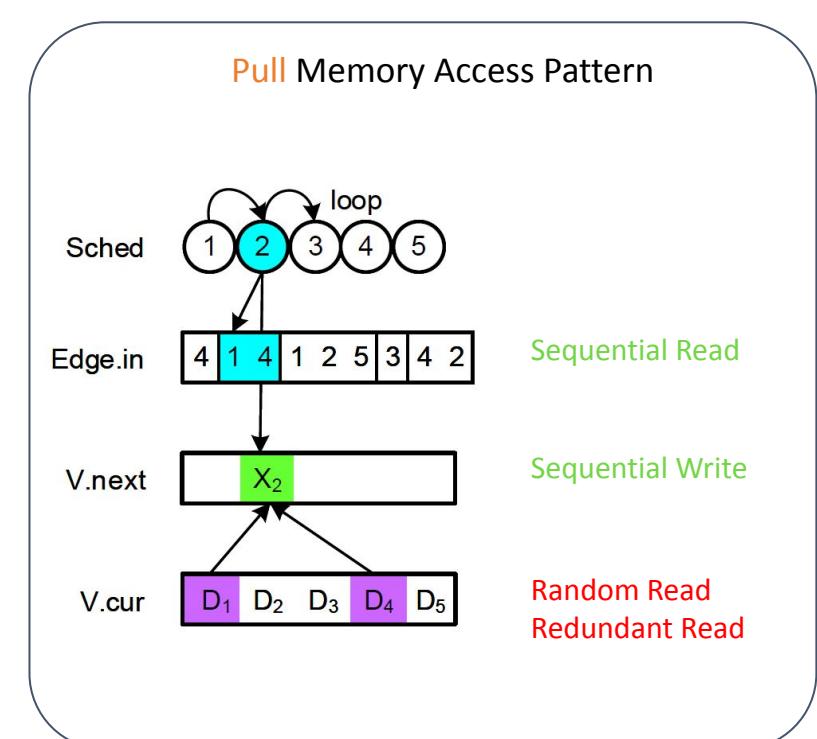
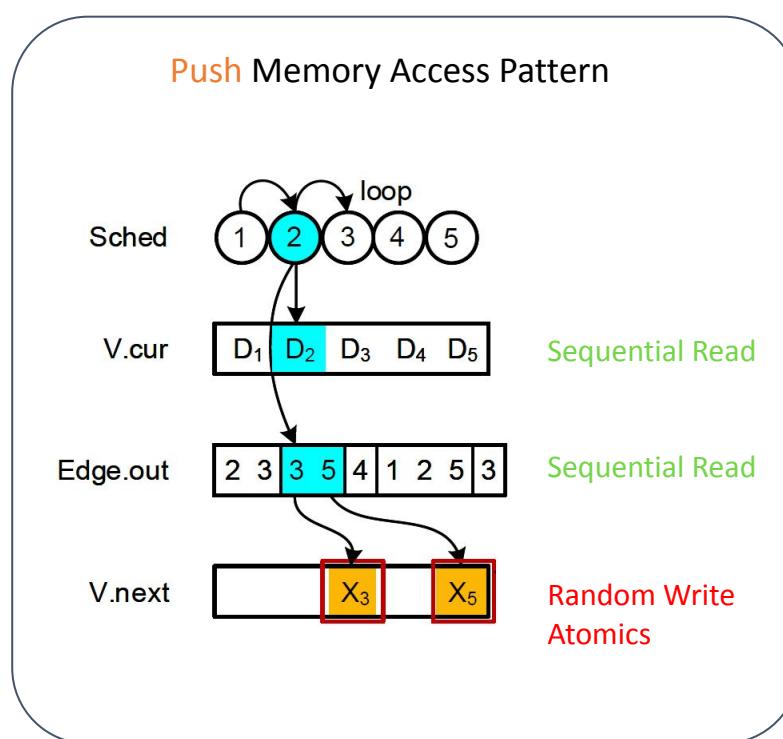
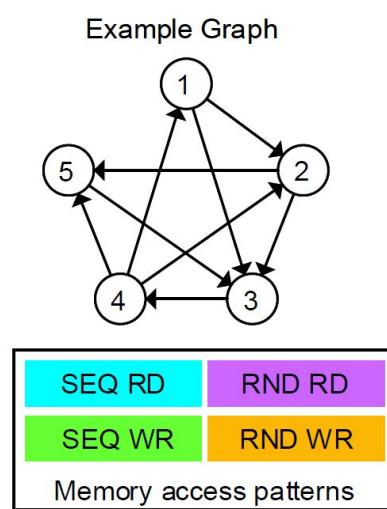


EECS 627 Final Presentation

Group 6 - An Event-Driven Graph Processing Accelerator

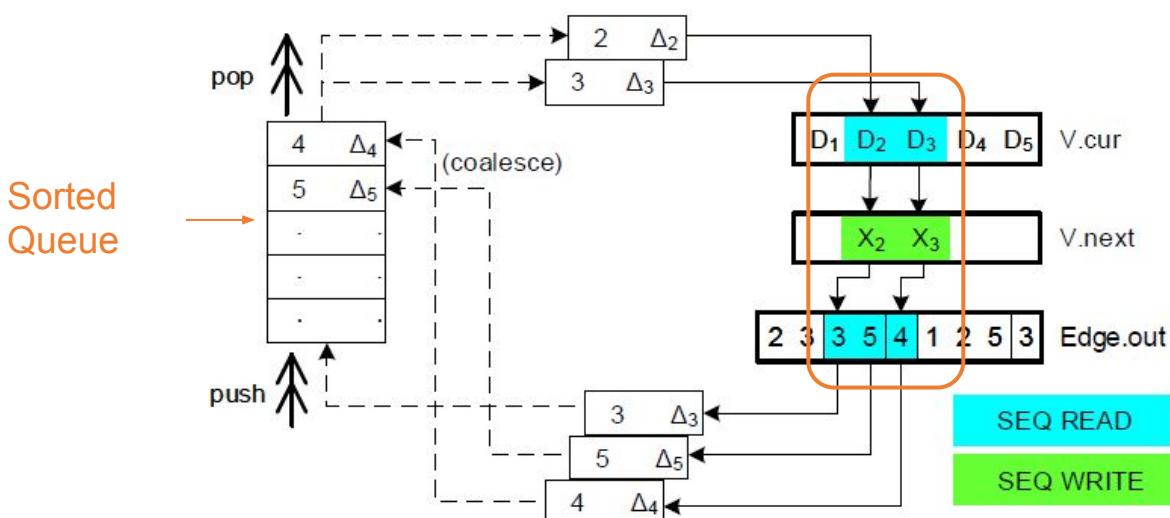
Motivation

- Vertex-centric Graph Processing Models either **Push/Pull** updates to/from neighbours
 - Both involves expensive **random memory access**



S. Rahman, N. Abu-Ghazaleh and R. Gupta, "GraphPulse: An Event-Driven Hardware Accelerator for Asynchronous Graph Processing," 2020 53rd Annual IEEE/ACM International Symposium on Microarchitecture (MICRO), Athens, Greece, 2020, pp. 908-921, doi: 10.1109/MICRO50266.2020.00078.

Random Access Reduction



Algorithmic insight:

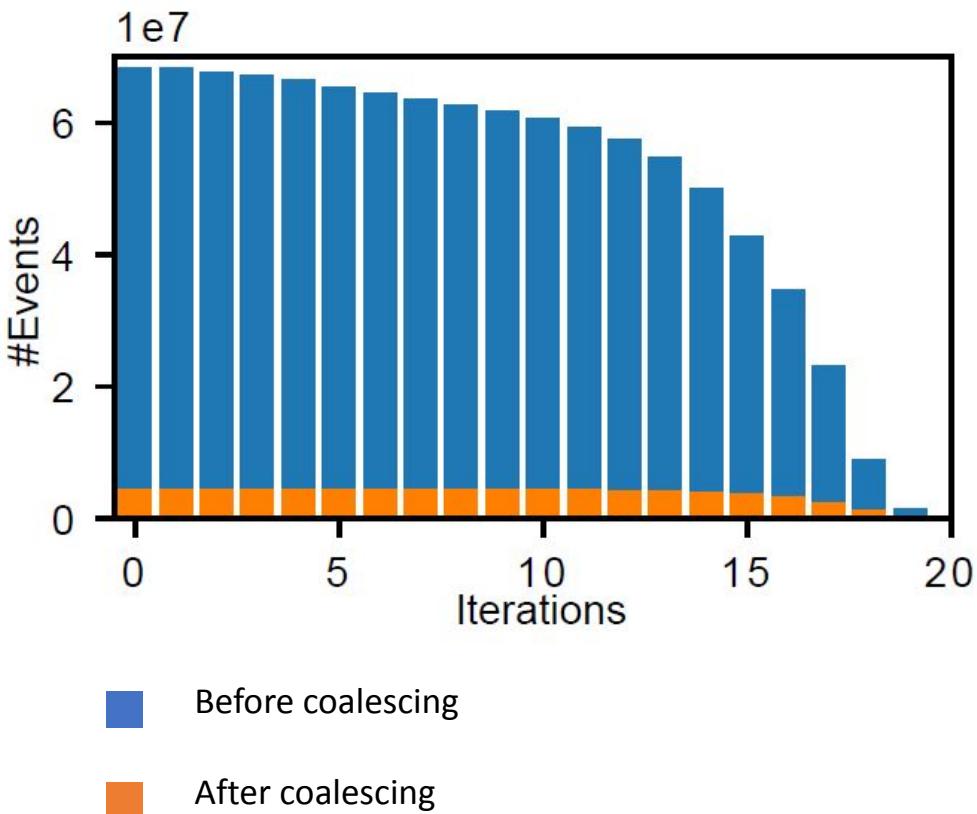
- Reduce function is **commutative** and **associative** for many algorithms
- Events targeting the same vertex can be processed in any order

Solution

- Sort the events with target vertex in the event queue
- Issue events with consecutive target vertices at the same time
- Convert random access to sequential access

S. Rahman, N. Abu-Ghazaleh and R. Gupta, "GraphPulse: An Event-Driven Hardware Accelerator for Asynchronous Graph Processing," 2020 53rd Annual IEEE/ACM International Symposium on Microarchitecture (MICRO), Athens, Greece, 2020, pp. 908-921, doi: 10.1109/MICRO50266.2020.00078.

Event Coalescing



Problems

- A huge number of events are generated in each iteration
- Massive storage of events
- Atomic operations

Algorithmic insight

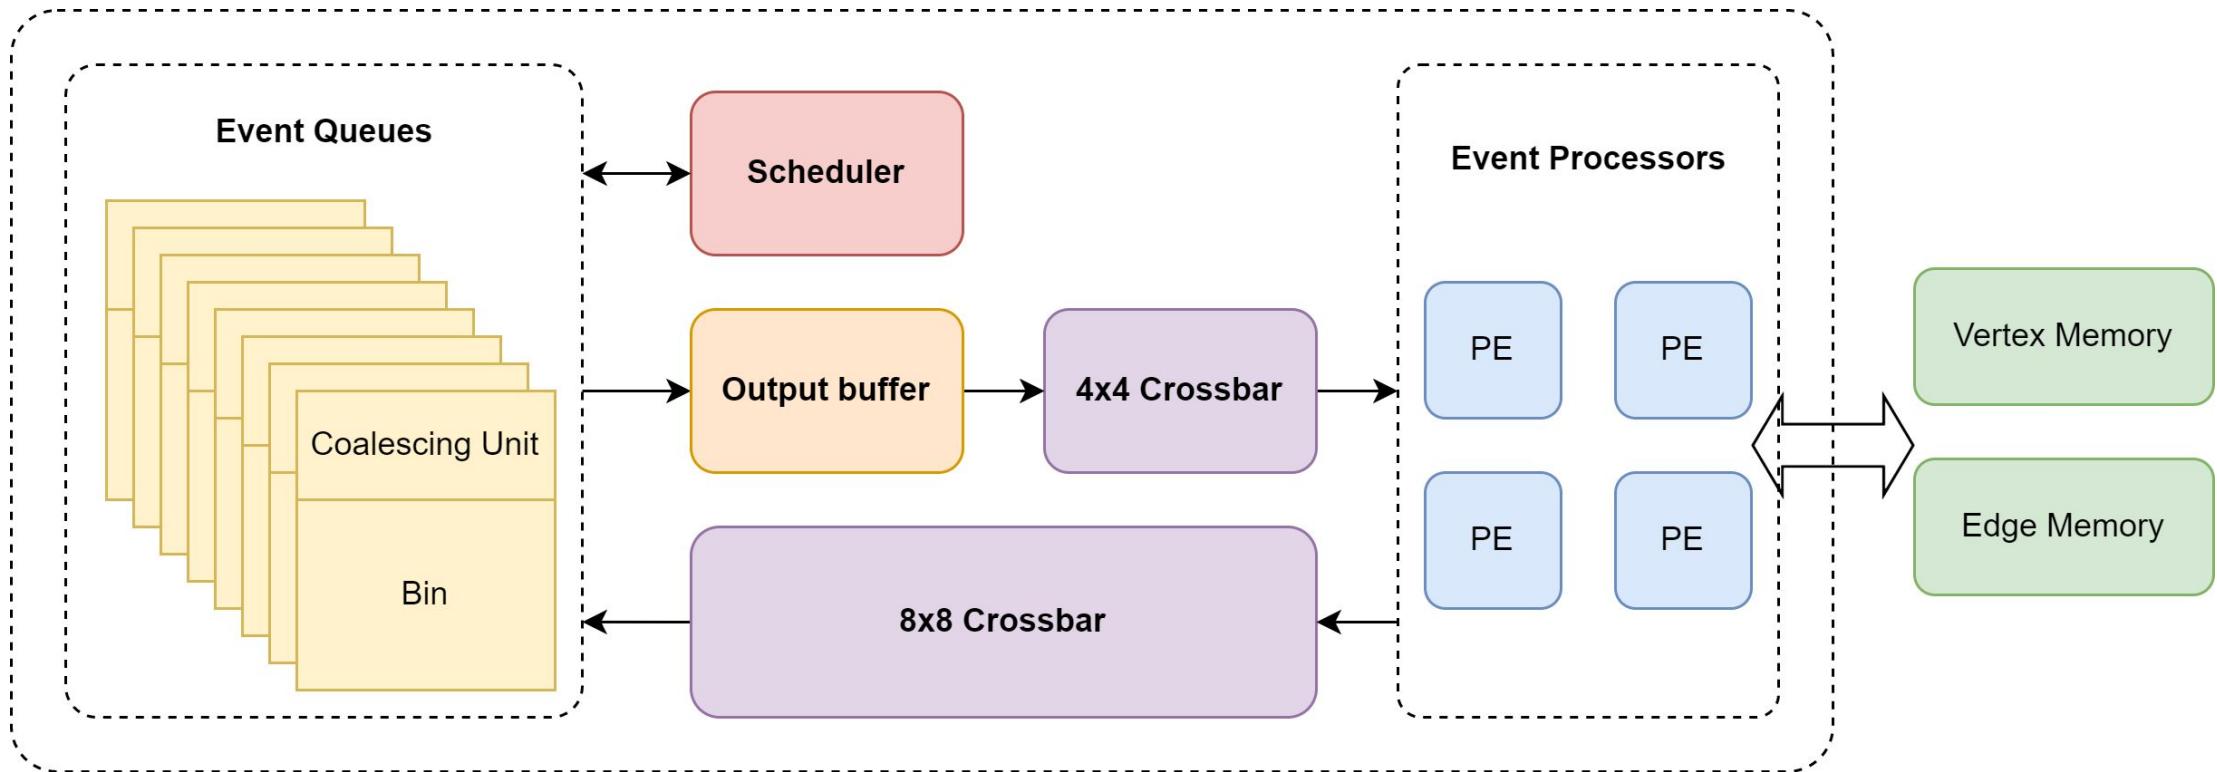
- Propagation function is **distributive** for many algorithms

Solution: Coalesces events targeting the same vertex using Reduce function in the event queue

- Less **storage** of events
- Less **workload** to be processed
- **Only one event targeting each vertex** at a time
- Faster convergence

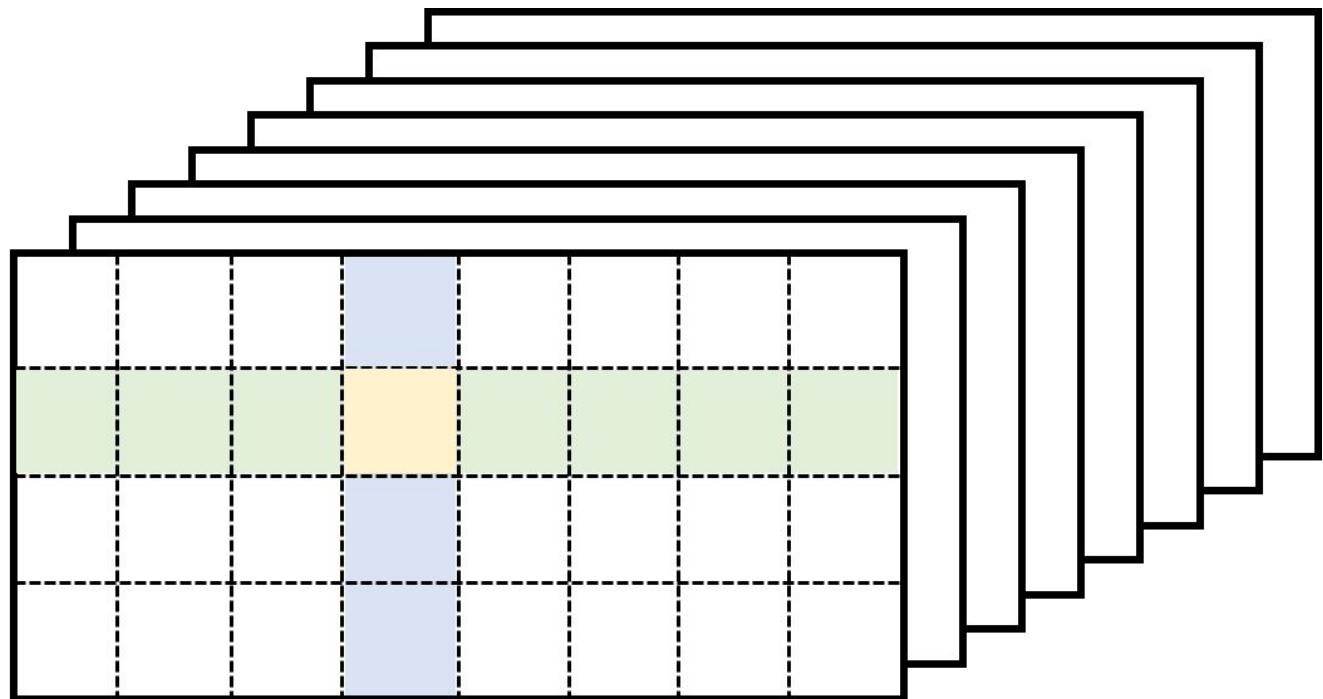
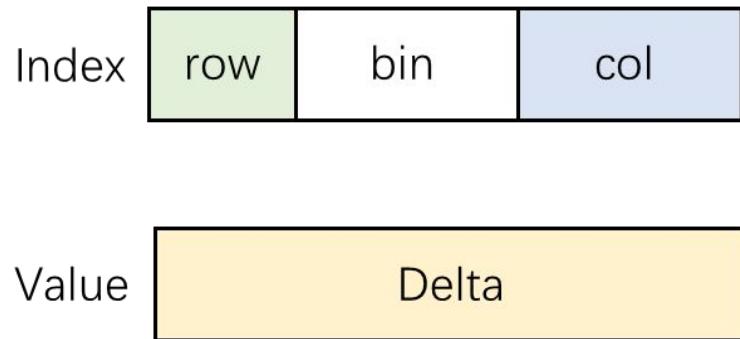
S. Rahman, N. Abu-Ghazaleh and R. Gupta, "GraphPulse: An Event-Driven Hardware Accelerator for Asynchronous Graph Processing," 2020 53rd Annual IEEE/ACM International Symposium on Microarchitecture (MICRO), Athens, Greece, 2020, pp. 908-921, doi: 10.1109/MICRO50266.2020.00078.

Project Overview



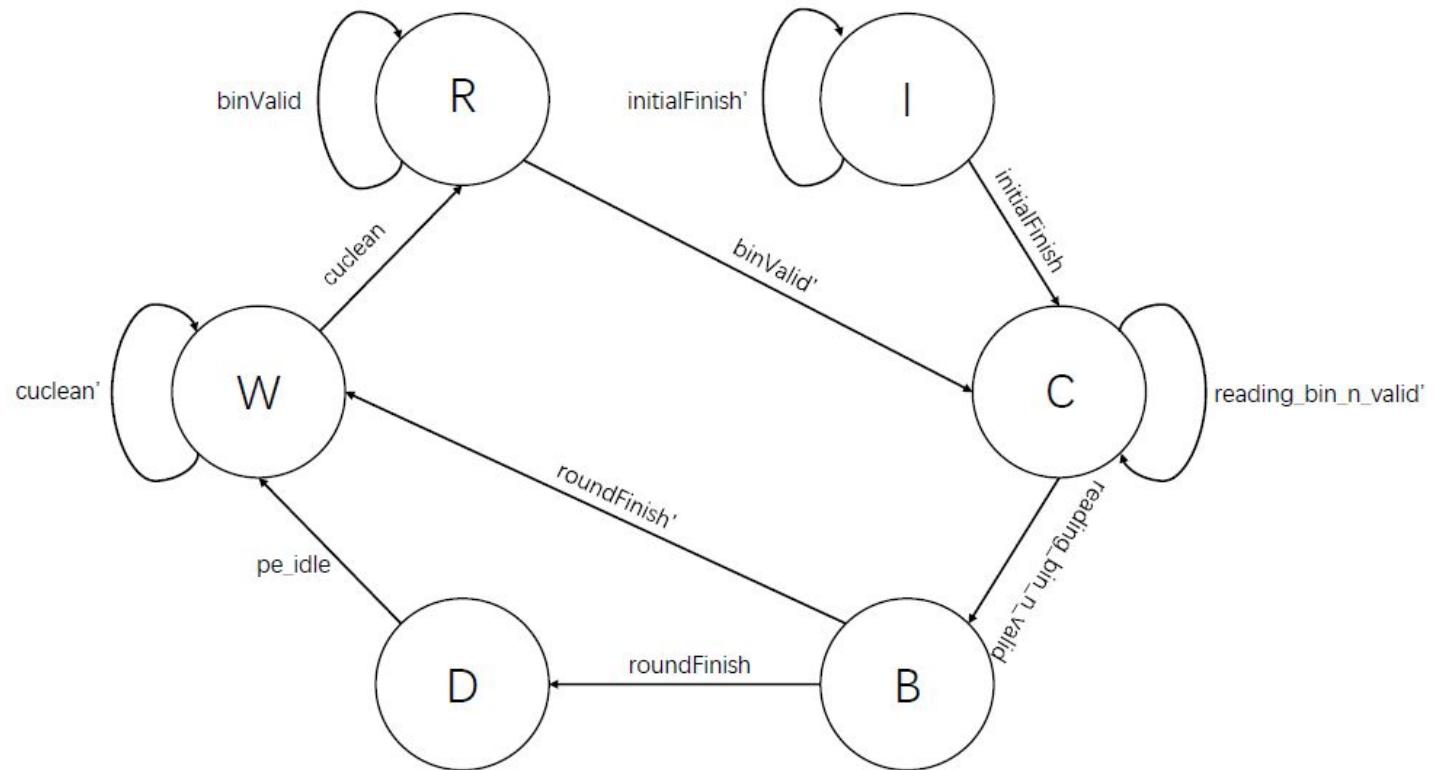
Event Queues

- Direct mapped event storage
 - 8 bin
 - 4 row
 - 8 column



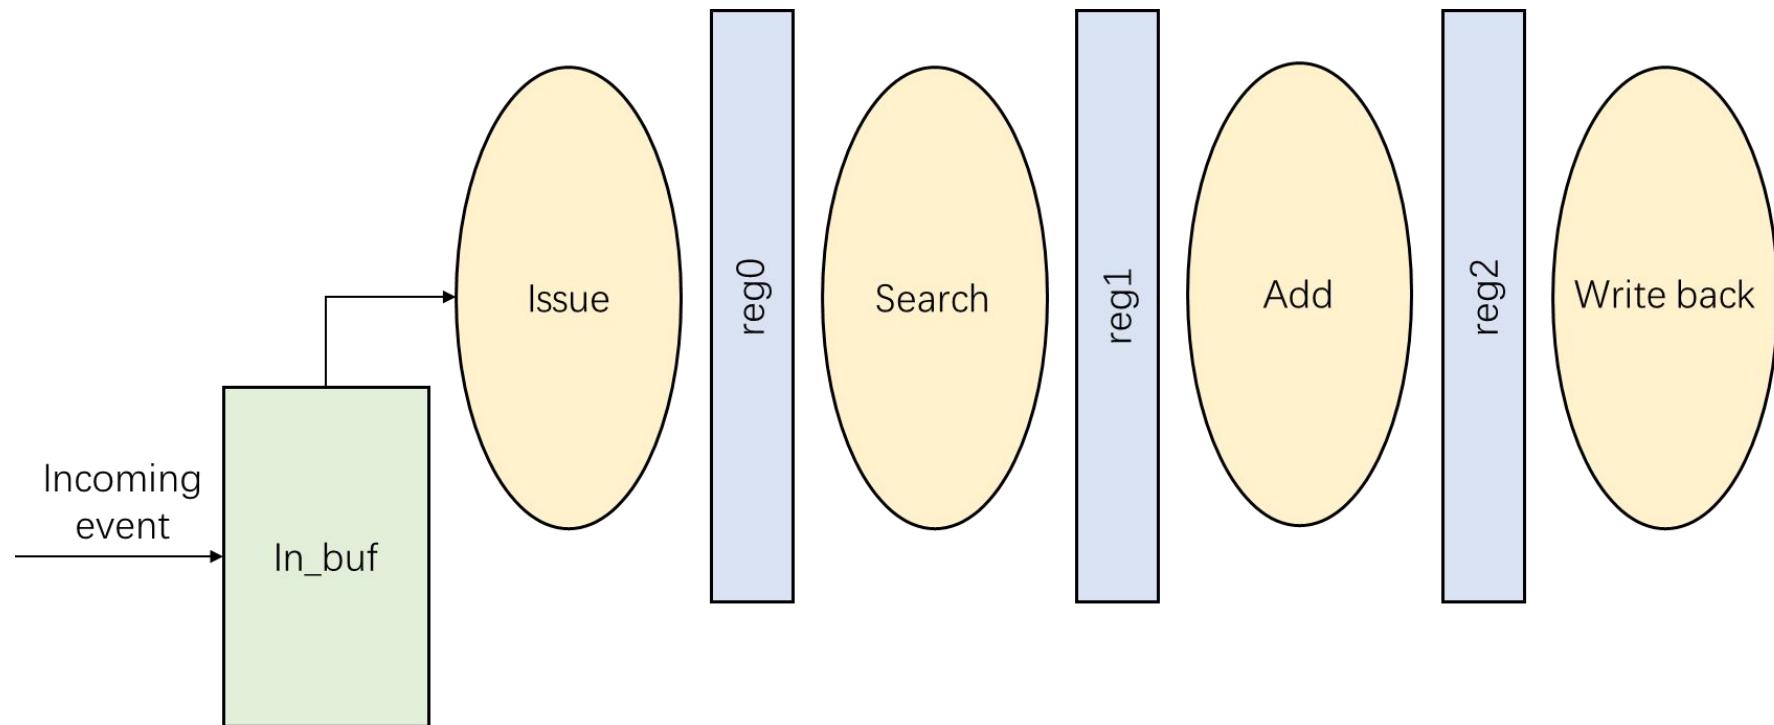
Scheduler FSM Modified for Atomicity

- 6-state FSM
 - **I**: initialization
 - **C**: start to select the next reading bin
 - **B**: stop taking new events from the input buffer into CU registers of the selected reading bin
 - **D**: wait until PE all idle when one bin is selected again
 - **W**: wait for CU to be clean
 - **R**: read enable



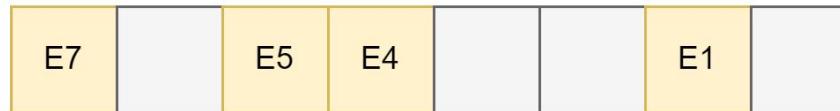
Coalescing Unit

- Pipelined coalescing with input buffer
- Combine multiple events to the same destination



Scheduler Output Buffer

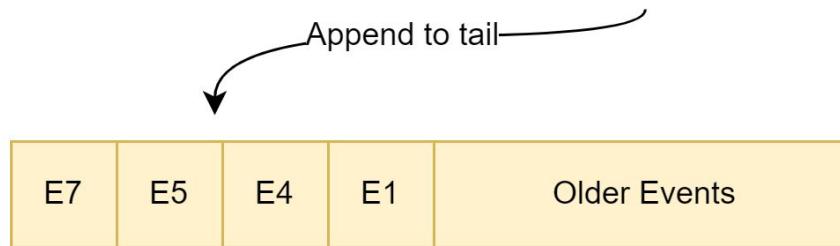
0. One whole row from Event Queues. Might be sparse.



1. Squeeze the bubbles (invalid entries) and store in a buffer

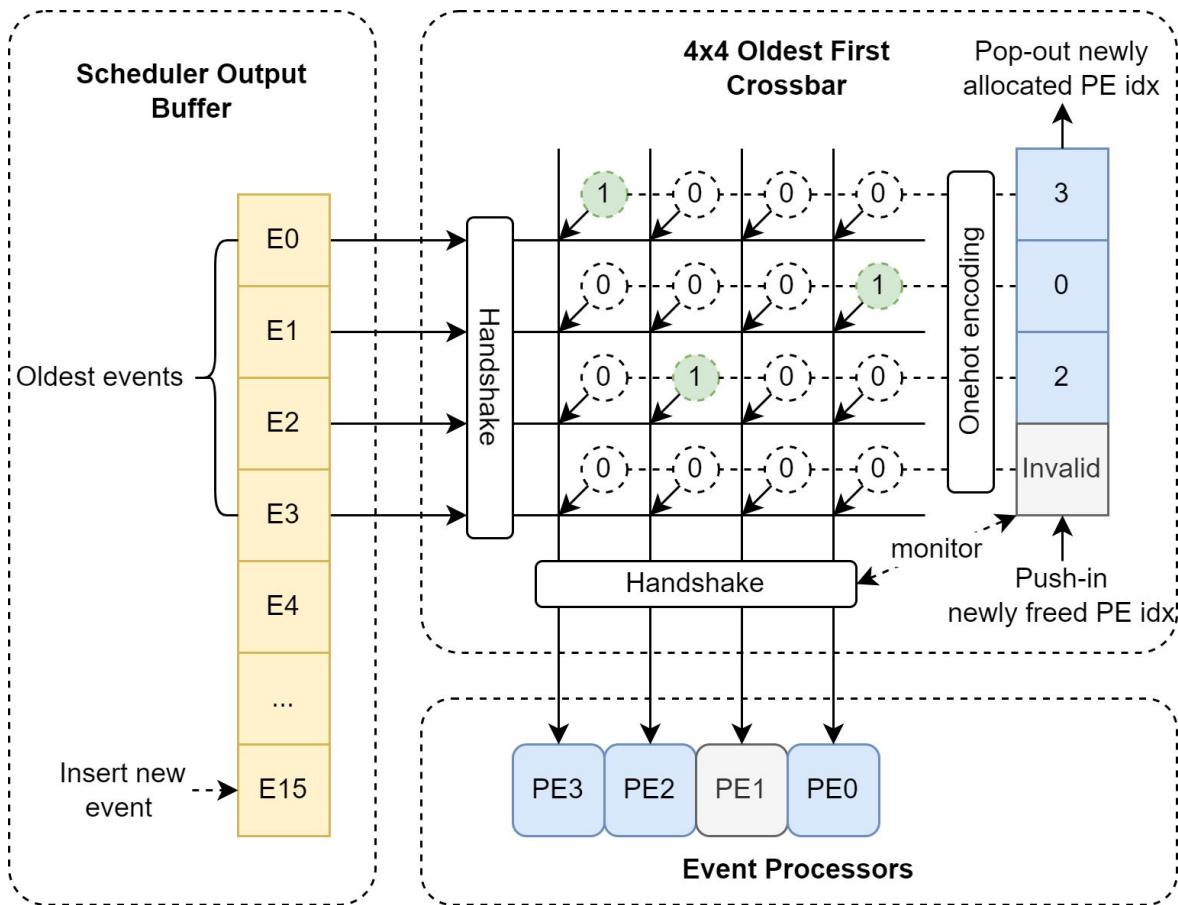


2. Append to the tail of a FIFO



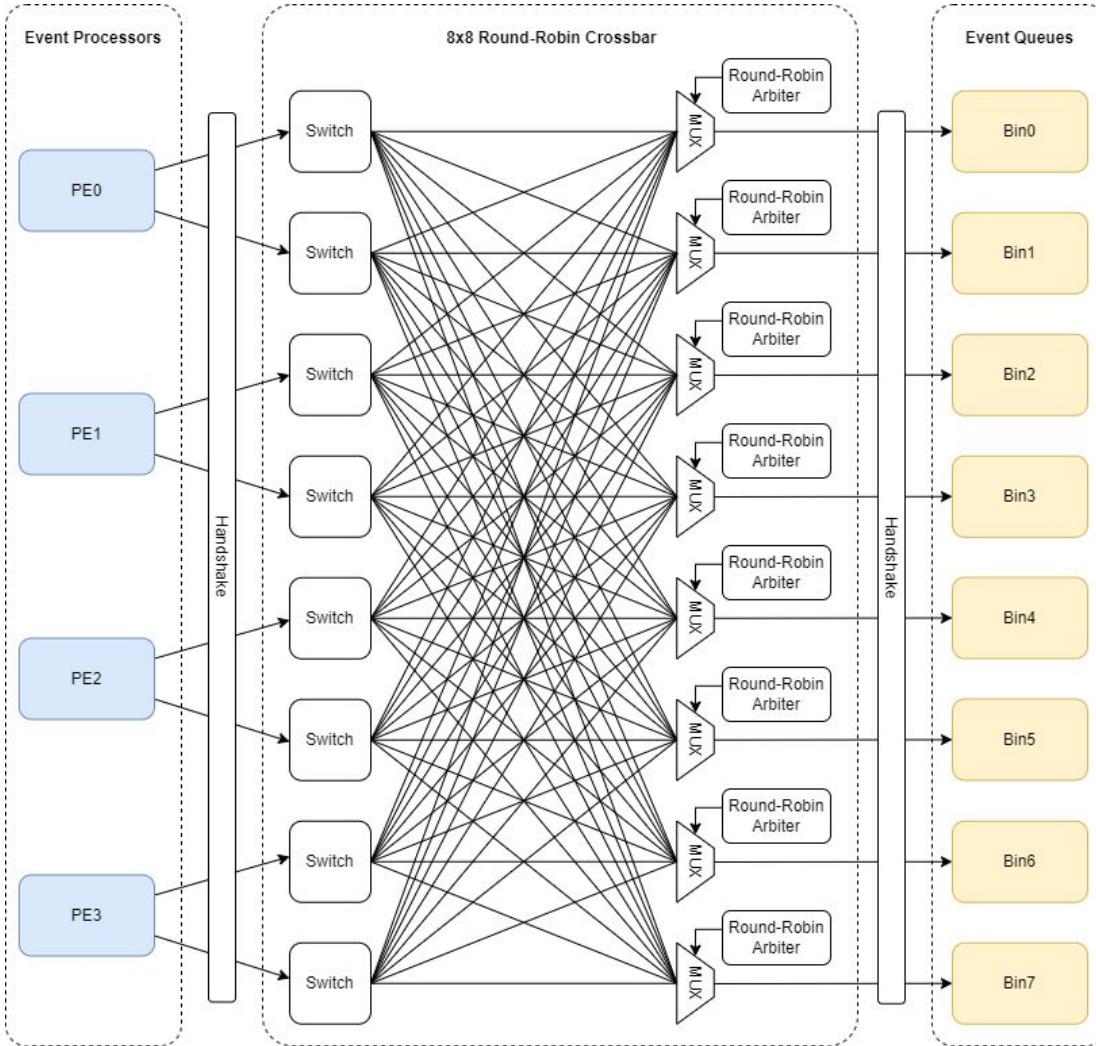
Crossbar to PE

Crossbar - from Scheduler to PEs



- Oldest events are guaranteed to be allocated first.
- Earliest freed PEs are guaranteed to be allocated first.

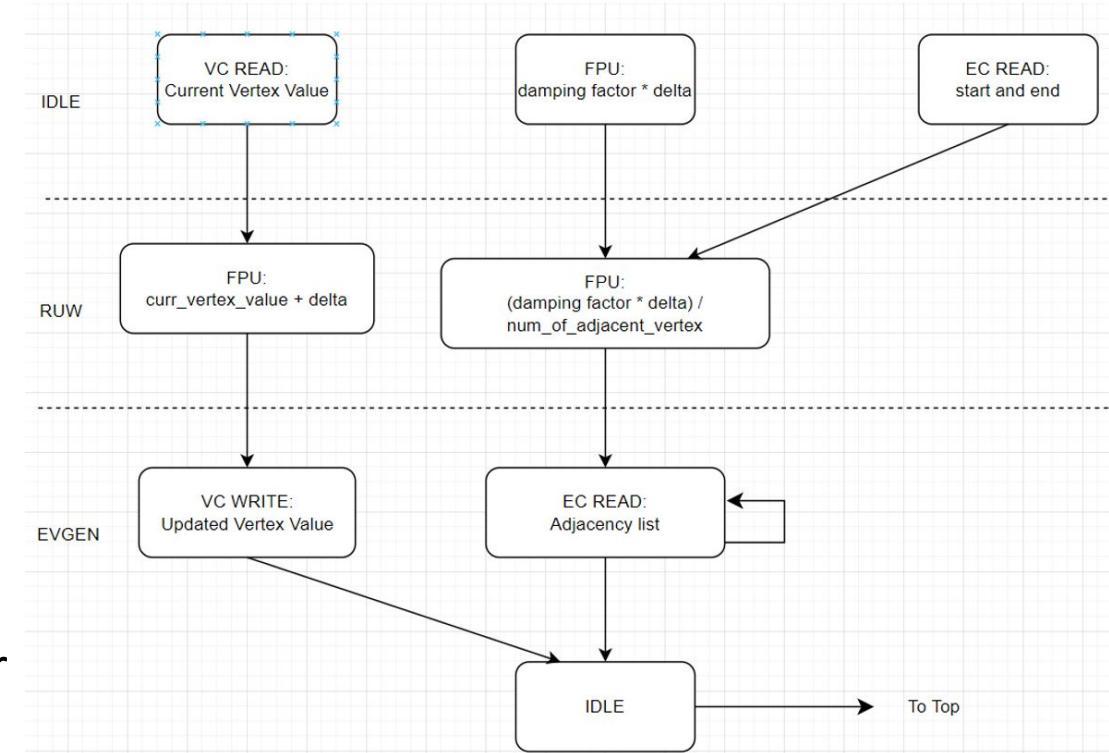
Crossbar - from PEs to Queues



- Switches direct the generated events to a bin according to vertex ID.
- Round-Robin arbiters guarantee fairness among PEs.

Event Processors

- Process existing events and generating new events (propagate)
- 4-state FSM
 - INIT: generating initial event sequence
 - IDLE: waiting for valid incoming event
 - RUW: calculating current event result
 - EVGEN: generating propagation events
- FPU
 - 4-stage pipelined as mentioned in paper
 - 16-bit half-precision (IEEE 754)
- 4 processors in our system

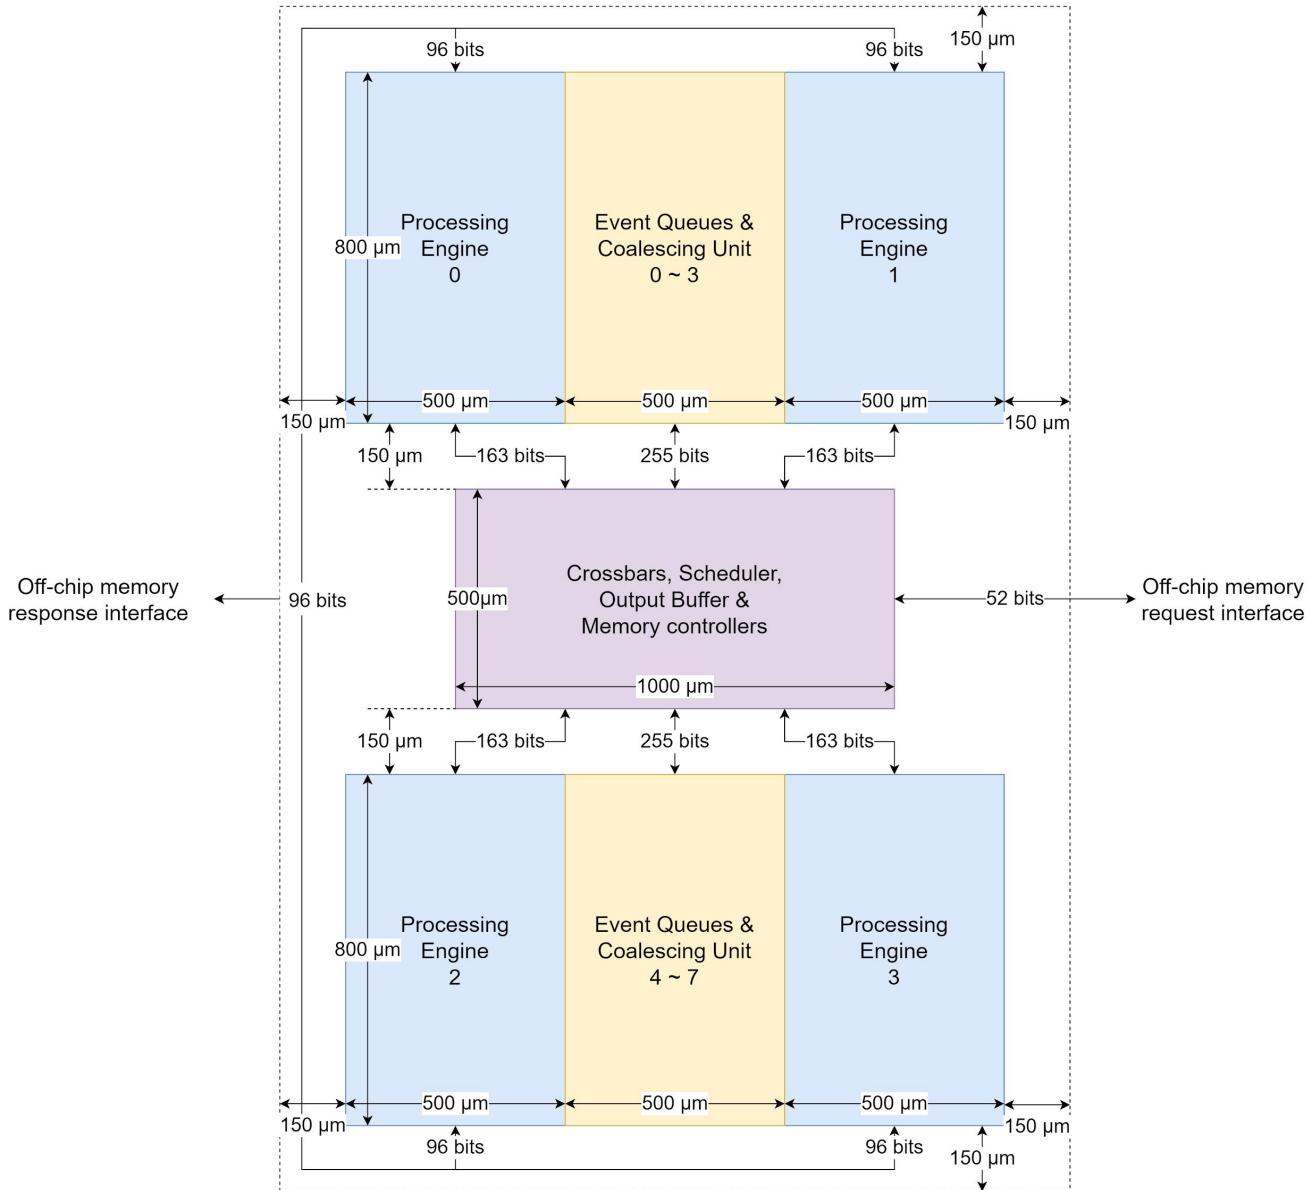


Memory

- No on-chip cache - off-chip memory models from EECS470
 - 64 bits per address
 - Configurable size
 - 5-cycle latency
- Two memories
 - One for vertex, one for edge/adjacency lists
 - Vertex: one 16-bit value per address
 - Edge:
 - CSR
 - up to four 16-bit row indices or eight 8-bit column indices per address

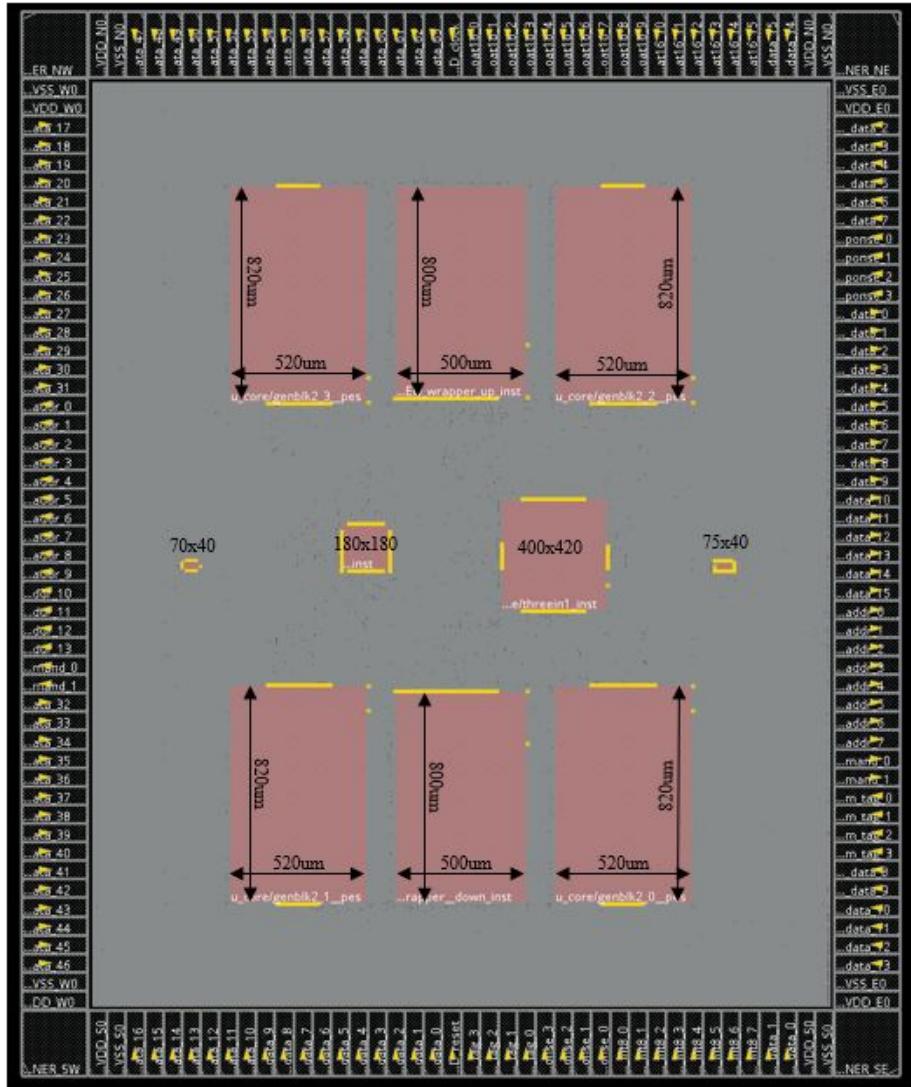
Floorplan

- EQ are divided into 2 groups.
- PEs are of similar distance to shared resources.
- Shared resources a group together to get better area efficiency.
- Gaps are reserved for top level routing.
- Total area estimation
 - Width: 1.8 mm
 - Height: 2.1 mm
 - Area: 3.78 mm^2



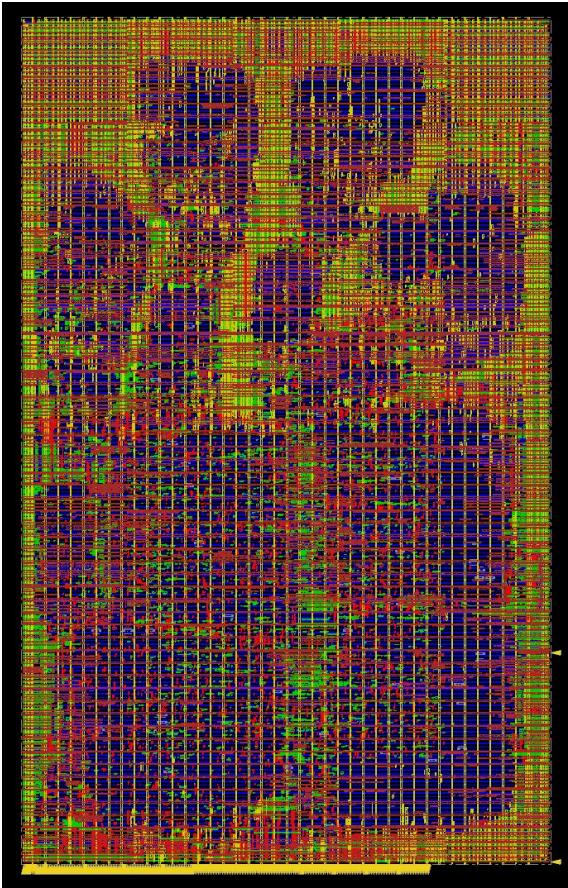
Floorplan

- To resolve DRC by increasing module density, the flattened center module had to be divided into subblocks
- For better routing to and from pads, modules are placed with more gaps in between
- Area with Pad ring:
 - Width: 3.3mm
 - Height: 4mm
 - Area: 13.2mm^2

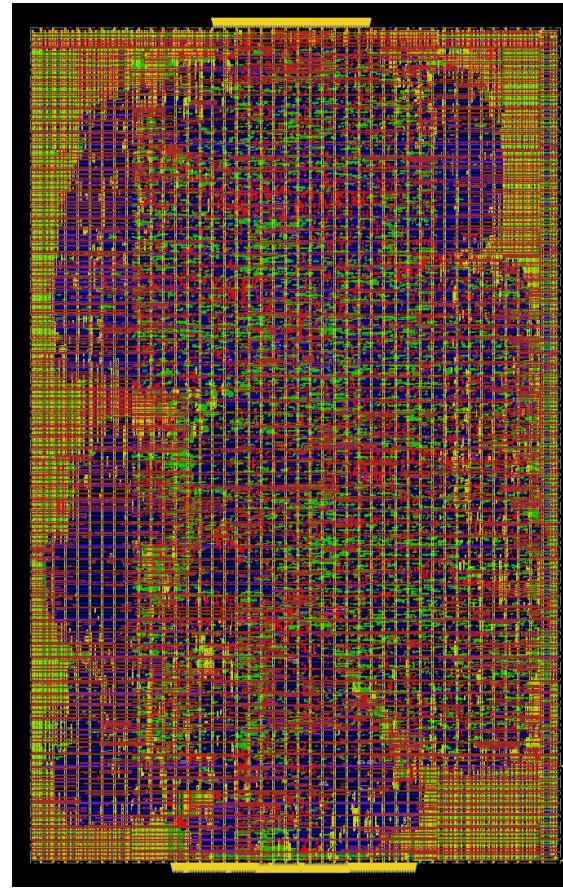


Module Level Layout

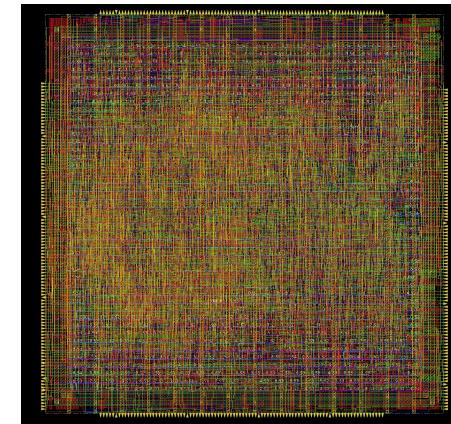
EQ (4 bins): 500x800



PE: 520x820

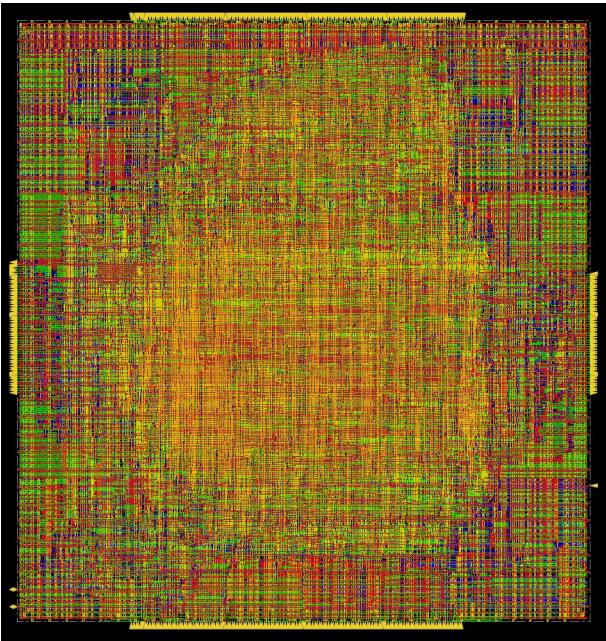


xbar: 180x180

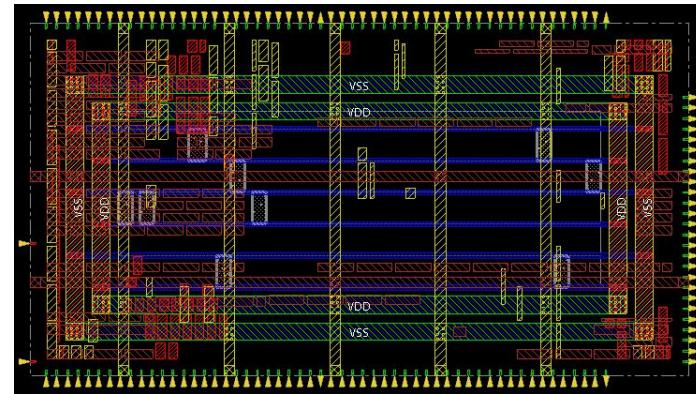


Module Level Layout

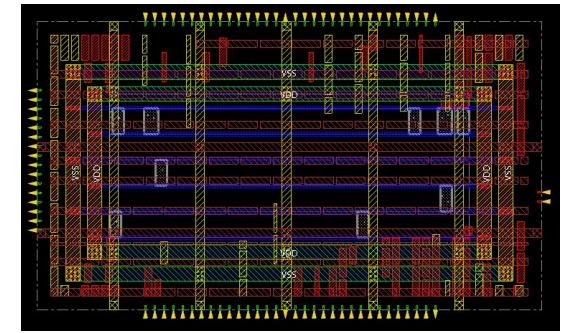
3in1: 400x420



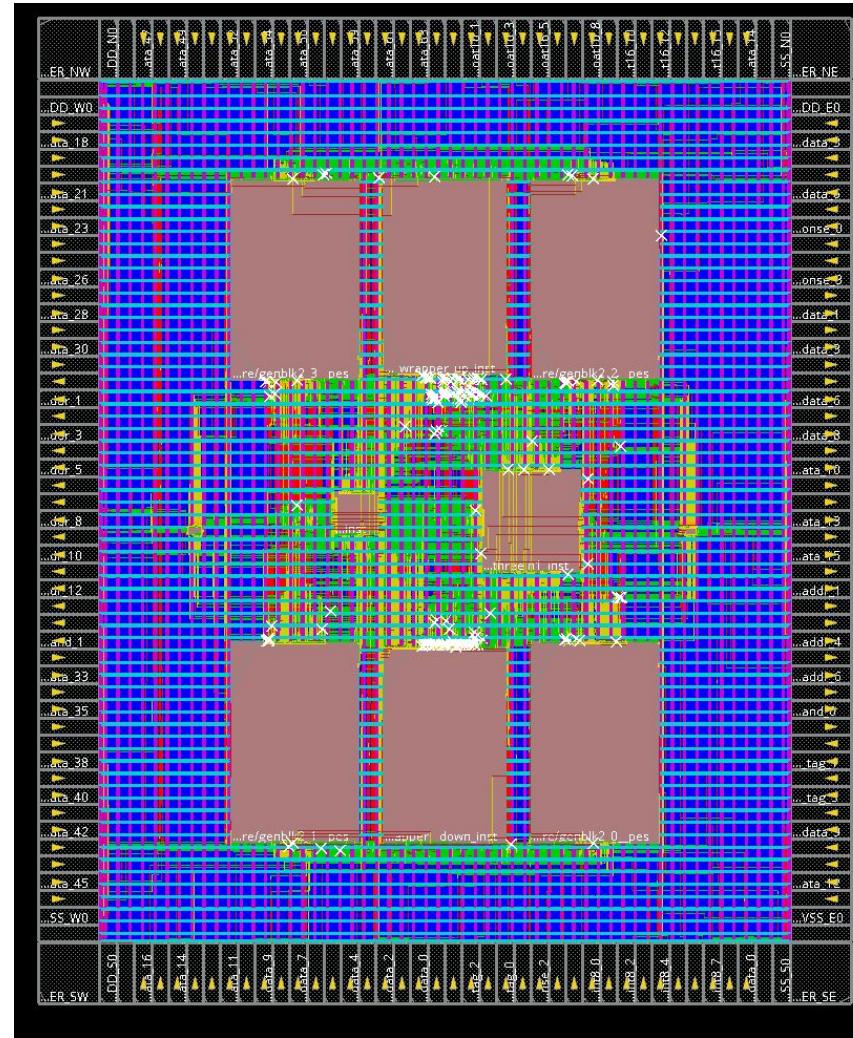
MC_VM: 75x40



MC_EM: 70x40



Top level Layout



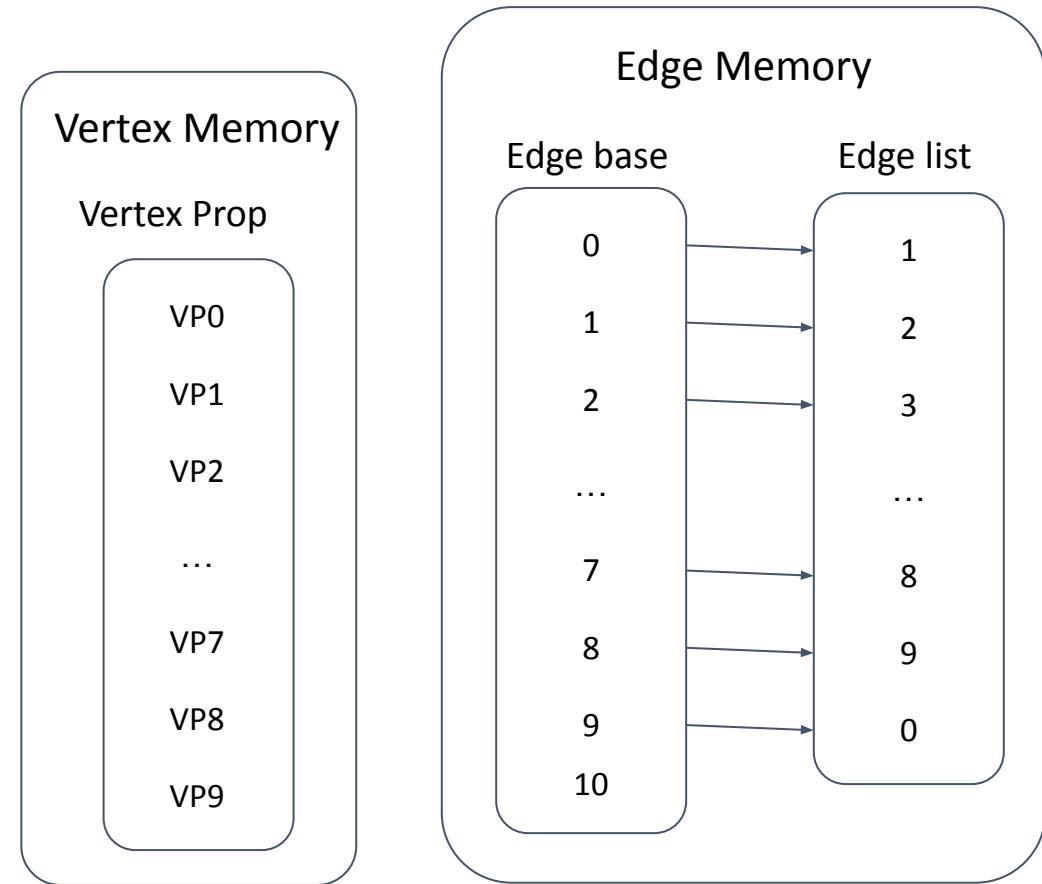
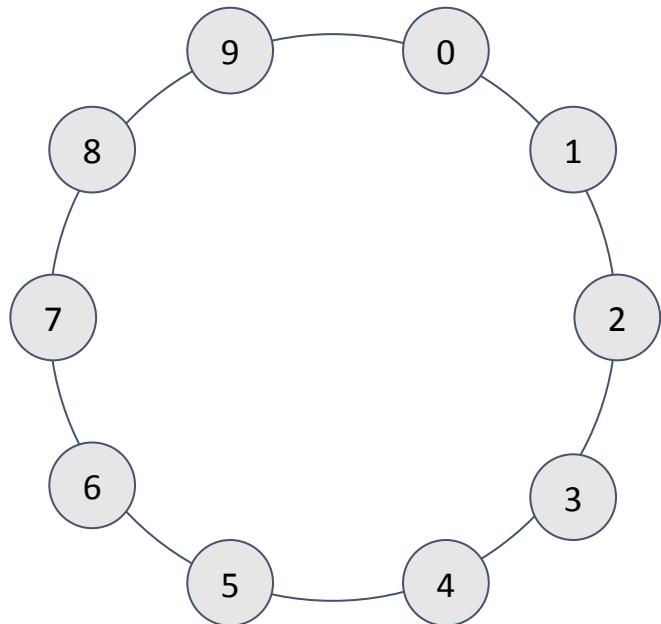
Golden Brick Methodology

- A cycle accurate golden brick (Verilog style) in Python
 - global variables in `io_port.py`
`signal, signal_n`
 - `one_clock()` method in each module class
`io_port.signal_n = ...`
 - `update()` in the main function
`io_port.signal = io_port.signal_n`
- Advantage: easy to transfer into Verilog code
- Disadvantage: inefficient in development

```
in main function:  
module0.one_clock()  
module1.one_clock()  
module2.one_clock()  
...  
update()
```

Simulation Vector

- PageRank on a 10-vertex graph in the shape of a ring



Pre-synthesis Verification

- Verification in SystemVerilog
- Compare the golden output with RTL output
 - All vertex values end up with 0.1

```
vertex mem [      0] = 2e61
vertex mem [      1] = 2e61
vertex mem [      2] = 2e61
vertex mem [      3] = 2e61
vertex mem [      4] = 2e61
vertex mem [      5] = 2e61
vertex mem [      6] = 2e61
vertex mem [      7] = 2e61
vertex mem [      8] = 2e61
vertex mem [      9] = 2e61
Time out @      5001
$finish called from file "testbench.sv", line 174.
$finish at simulation time      250275000
V C S   S i m u l a t i o n   R e p o r t
Time: 250275000 ps
```

Post-synthesis Verification

- **Problem**
 - Synthesized under 10ns without negative slacks
 - Many X values in netlist
- **Cause - X propagation in netlist**
 - Some values were assigned X in RTL
 - The X values propagated through the netlist
 - Unable to detect in RTL
- With modification on RTL code, the top level netlist pass the simulation under the clock cycle of 22ns

Design Metrics

	Area(um^2)	Density	T_clk(ns)	Setup	Hold
Xbar	180x180	95.97%	10	5.001	0.061
EQ	500x800	52.48%	10	0.289	0.066
MC_EM	70x40	75.84%	10	8.338	0.179
MC_VM	75x40	96.93%	10	8.216	0.178
PE	520x820	61.16%	10	0.431	0.105
3in1	400x420	40.10%	10	0.431	0.105
Top level	3300x4000		11	4.096	0.211

clock uncertainty: 0.1ns

Design Metrics

	Switch Power	Int Power	Leak Power	Total Power
Xbar	0.29	0.762	5.89e+05	1.052
EQ	1.055	5.044	5.36e+06	6.105
MC_EM	4.24e-02	2.87e-02	1.82e+04	7.11e-02
MC_VM	6.06e-02	3.22e-02	2.40e+04	9.28e-02
PE	6.426	5.726	1.31e+07	12.165
3in1	0.571	1.362	1.13e+06	1.934
Top level	0.728	19.549	4.75e+07	20.325

Dynamic Power Units = 1mW, Leakage Power Units = 1pW. Post-synthesis.

Final Status

- All sub-modules' APR are DRC/LVS clean
- Post-synthesis verification passed at clock period of 22ns
- Top Level Integration:
 - Pad ring, floorplan and routing are proved to be feasible
 - Power grid generation is not successful. Possible reason:
 - Sub-modules are configured to use M1~M6 for routing
 - Sub-module power-grids are mistakenly placed on M4 & M5
 - The top-level APR cannot correctly connect top-level power grid (M7 & M8) to sub-module power.
 - Not enough time to rerun all the sub-module APR.

Looking back...

- Graph computation is fun.
- A novel architecture without published implementation work.
- No open-source simulation code. Have to construct golden brick from scratch.
- Complicated RTL design and verification, due to its asynchronous nature.
- DRC/LVS clean sub-modules does not necessarily lead to top-level success.
- Should have paid more attention to power grids.